

Cluster contains multiple nodes but there will be only one master node.

kubectl version: to get version of kubernetes

POD

kubectl run <pod_name> --image=<image_name> -> it will run a new pod with image.

kubectl run nginx --image=nginx -> it will run a new pod called nginx with image as nginx.

kubectl get pods -> to get all the pods

kubectl get pods -o wide -> to get more details for the pods

kubectl describe pod <pod-name> -> it will give us the detailed info about the specific pod

Kubernetes yaml file must contain these fields -> apiVersion, kind, metadata, spec

kind	version
Pod	v1
Service	v1
ReplicationController(deprecated)	v1
ReplicaSet	apps/v1
Deployment	apps/v1

nginx-pod.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  tier: frontend
spec:
  containers:
    - name: nginx
      image: nginx
```

kubectl apply/create -f nginx-pod.yaml -> to create a pod with this yaml file.

kubectl describe pod nginx -> to give more description of the nginx pod

kubectl edit pod nginx -> it will open a vi editor where we can change the pod definition file also this is a in memory pod definition file which is maintained by Kubernetes.

kubectl delete pod nginx -> to delete the nginx pod

kubectl delete --all pods -> delete all the pods.

kubectl delete all --all -n {namespace} -> delete all resources in namespace.

kubectrl run redis --image=redis --dry-run=client -o yaml > redis-pod.yaml -> It will not create any pod rather it's a imperative style of writing definition file where a pod definition file will created with the necessary fields.

ReplicaSet

ReplicaSet is a group of same pods where we can scale in(reduce) or scale out(increase) the number of the pods.

kubectrl create/apply -f <replicaset-definition.yaml> -> it will create a replicaset from the definition file.

kubectrl get replicaset -> to get all the replicaset in the default namespace

kubectrl describe replicaset -> to get all the replicaset in the default namespace

kubectrl delete replicaset <replicaset-name> -> to delete the replicaset

nginx-replicaset.yaml

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: nginx-replicaset
  labels:
    type: frontend
spec:
  selector:
    matchLabels:
      name: nginx-pod
      type: frontend
  replicas: 1
  template:
    metadata:
      labels:
        name: nginx-pod
        type: frontend
    spec:
      containers:
        - name: nginx
          image: nginx
```

labels under template section and mathLabels under selector should be same. That is how replicate set identifies the pod and controls the number of the pods. If we try to delete any pod or anyhow any pod got crashed, then Kubernetes automatically brings another pod in.

Scale commands ->

kubectrl replace -f nginx-replicaset.yaml -> replace the previous nginx-replicaset with the current replicas given in the definition file

kubectrl scale --replicas=6 -f nginx-replicaset.yaml -> it will override the replicas given in the yaml

kubectl edit replicaset nginx-replicaset -> it will open an editor and show the current configuration of the replicaset then we can easily scale out.

kubectl scale --replicas=6 replicaset nginx-replicaset -> it will scale out an existing replicaset

Deployment

Kubernetes deployment create creates one deployment kind of object.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx-deployment
spec:
  replicas: 2
  selector:
    matchLabels:
      app: nginx-pod
  template:
    metadata:
      name: nginx-pod
      labels:
        app: nginx-pod
    spec:
      containers:
        - name: nginx
          image: nginx
```

kubectl get deployments -> It will fetch all the deployments.

kubectl describe deployments <deployment-name> -> It will give us the description of the specific deployment

kubectl get all -> To get the resources like pod, replicaset, services, deployments

kubectl create deployment http-frontend --image=httpd:2.4.alpine -> it will create one deployment named ad http-frontend and its image will be httpd:2.4.alpine and count of the pod will be 1.

kubectl scale deployment --replicas=3 httpd-frontend -> it will scale the deployment and change the number of the pods

kubectl create -f nginx-deployment.yaml --record -> It will create a deployment and also record all the changes of the deployment for rollout history.

If we want to revert to previous version, we must add this record flag

Update and roll back ->

kubectl rollout status deployment/nginx-deployment -> it will give the current rollout status of the deployment. When we change the replicas or the image at that time it will give the information.

kubectl rollout history deployment/nginx-deployment -> It will give us all the rollout history of the Kubernetes deployment rollout

For rollout Kubernetes has 2 types of strategy

1. Recreate strategy
2. Rolling strategy

Kubernetes deployment object creates another replicaset and then it will add the pods into the replicaset

kubectl rollout undo deployment/nginx-deployment -> It will revert the latest changes back to the previous version.

Changes to the existing deployment:

kubectl set image deployment nginx-deployment nginx:old-image=nginx:new-image --record -> It will change the image of the deployment of the nginx deployment and will record it

kubectl edit deployment nginx-deployment -- -> it will open vi editor and open the current configuration of the nginx-deployment

Service

For accessing the pods from the outside of the container we need service.

There are three types of services.

1. NodePort
2. ClusterIP
3. LoadBalancer

NodePort can range from 30000 to 32767. The work of the NodePort is to listen to a particular port and forward it to another node.

```
#nginx-service.yaml
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
spec:
  type: NodePort
  ports:
    - port: 80
      targetPort: 80      #where the pod will listen (for nginx its 80)
      nodePort: 30008     #In this port the service will be accesible
  selector:
    app: nginx-pod       #to connect with specific pods via pod's label
```

kubectl create -f nginx-deployment.yaml -> for creating nginx-service with nginx-deployment.yaml

kubectl get services -> To get all the services

kubectl describe service <service_name> -> To get details of the specific service

minikube service nginx-service --url -> It will print the service url

ClusterIP: Internal communication of pods. Service definition is almost same as the NodePort. Here type is **ClusterIP**. **TargetPort** is where the backend is exposed, and **Port** is where the service is exposed. Type ClusterIp is the default service type.

LoadBalancer: With the NodePort service we can make external facing application available on the port of the worker nodes.

Let's say we have four cluster and one each server there are one frontend app deployed. With NodePort we can make external traffic forwarded to frontend pod but again for that we will have 4 urls for 4 services.

So, need to have a loadbalancer here. We can have an external VM where nginx is deployed and then it will loadbalance 4 urls but it will be a complicated thing to manage.

So we can use the inbuilt loadbalancer of different cloud platforms like Azure,GCP or AWS.

kubectl describe service <service-name> -> if there is no endpoint then the service is not attached to any pod.

minikube service <service-name> --url -> to get the url of the NodePort service.

Imperative style of creating service:

kubectl expose deployment nginx-deployment --name=nginx-service --target-port=80 --port=80 --type=NodePort -> It will create a service named as nginx-service of type NodePort with specific port

and targetPort and it will match labels of the deployment of nginx-deplpyment and NodePort will be assigned randomly in the range of 30000 to 32767

kubectl expose deployment nginx-deployment --name=nginx-service --target-port=80 --port=80 --type=NodePort --dry-run=client -o yaml > nginx-service.yaml -> It will do just the same as previous just that it will save all the configurations in the nginx-service.yaml

If we have a connection like this:

1. Voting-app -> frontend app for gathering the votes which will save the votes in redis
2. Result-app -> frontend app for showing the votes which will fetch votes from postgres
3. Redis -> save the votes from in a in memory store
4. Postgres -> save the votes in relational db
5. Worker-app -> It will constantly fetch the vote count from redis and constantly update the vote count in postgres

So, our setup will be like this:

1. Deploy the pods/deployment
2. Create the ClusterIP service for Redis and postgres
3. Create the NodePort service for Voting-app and Result-app

NameSpace

NameSpace is a way to segregate different resources like dev,qa,prod etc. Default namespace is default.

```
#dev-namespace.yaml
apiVersion: v1
kind: Namespace
metadata:
  name: dev
```

Kubectl create -f dev-namespace.yaml -> to create a namespace

kubectl get pods --namespace=dev -> To get all the pod inside dev namespace

If we have to set the dev namespace permanently then we can keep it inside the **KubeConfig**

Kubectl config set-context \$(kubectl config current-context) --namespace=dev

To limit the resources using in a specific namespace we can use resource quota

kubectl get ns -> to get all the namespaces

kubectl get ns --no-headers | wc -l -> To get count of the namespaces

In pod-definition.yaml in metadata we can add namespace to mention the namespace where the pod will be deployed

```
# nginx-pod.yaml
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
  namespace: dev
  labels:
    app: nginx-pod
spec:
  containers:
    - name: nginx
      image: nginx
```

kubect1 -n dev get pods --no-header -> to get the pods in dev namespace

In which namespace the nginx pod is deployed?

Ans: **kubect1 get pods --all-namespaces | grep nginx**

In the same namespace we can connect to another pod via pod name but to connect with other resources in another namespace we have to maintain the proper format.

Resource-name.namespace.resource-type.domain

example: **db-service.dev.svc.cluster.local**

db-service is the name of the resource, dev is the namespace, svc is the resource type, cluster.local is the domain

Some of the imperative style command

--dry-run -> by default all the command is run with this --dry-run. As soon as the command is run the resource will be created.

--dry-run=client -> It will not create the resource rather it will check the whole command and tell us that the command is correct or not.

-o yaml > resource-definition.yaml -> this will create the resource definition in an yaml format

kubect1 create deployment nginx --image=nginx --dry-run=client -o yaml > nginx-deployment.yaml -> It will first check the command is correct or not. If it is correct then will create one nginx-deployment file with the configuration given

kubect1 expose pod redis --port=6379 --name=redis-service --dry-run=client -o yaml > redis-service.yaml -> it will create a redis-service yaml with the configuration given

kubect1 create service clusterip redis-service --tcp=6379:6379 --dry-run=client -o yaml > redis-service.yaml -> it will create a dry run of the redis service and save it in redis-service.yaml

kubect1 expose pod nginx-pod --port=80 --name=nginx-service --type=NodePort --dry-run=client -o yaml > nginx-service.yaml -> It will expose pod named as nginx pod type of NodePort4

kubectkl create service nodeport nginx-service --tcp=80:80 --node-port=30080 --dry-run=client -o yaml > nginx-service.yaml -> It will create a service of nodeport with port 80 and tagetport 80 with nodeport 30080

kubectkl expose command automatically use the pods labels as the selctors but we can not specify the nodeport. We must add that in the definition file then we can add the nodeport.

Kubectkl create service command will not use the pod labels as selectors instead it will assume selectors as **app: service-name** and we can not pass selector in the definition file.

Commands and args

In docker we have **ENTRYPOINT** and **CMD** for giving command line arguments, but we can override that using **--entrypoint** and the **extra parameters** passed in the docker run command. Same way we can override the existing command in the pod definition file with **command** and **args**. **ENTRYPOINT** will associated with **command** and **CMD** will be associated with **args**.

Let's say we have a dockerfile like this

```
#ubuntu-sleeper dockerfile
FROM UBUNTU
ENTRYPOINT ["sleep"]
CMD ["10"]
```

```
name: ubuntu-sleeper
image: ubuntu-sleeper
command: ["sleeper2.0"]
args: ["100"]
```

The right-side pod definition file is same as this docker run command:

docker run --entry-point=sleeper2.0 ubuntu-sleeper 100

Copy the content of existing pod into a yaml -> **kubectkl get pod <pod-name> -o yaml > pod-definition.yaml**

Environment variables

We can pass environment variables in these 3 ways.

1. Using env
2. Using config map key valueFrom -> configMapKeyRef
3. Using secret key valueFrom -> secretKeyRef

```
spec:
  name: postgres
  image: postgres
  env:
    - name: POSTGRES_PASSWORD
      value: MY_SECRET_PASSWORD
```

```
spec:
  name: postgres
  image: postgres
  env:
    - name: POSTGRES_PASSWORD
      valueFrom:
        configMapKeyRef:
          name: posgtres-config
          key: password
```

```
spec:
  name: postgres
  image: postgres
  env:
    - name: POSTGRES_PASSWORD
      valueFrom:
        secretKeyRef:
          name: postgres-secret
          key: password
```


We can either use **configMapKeyRef** where we can pass keys, or we can use **configMapRef** where we can directly pass the **configMap**. Same goes with Secrets. We can either use **secretKeyRef** or **secretRef**.

We can create configMap with these 3 ways.

1. By passing all the values in command.
2. By passing the properties file or yaml file in the command
3. By using the definition file.

kubectl create configmap app-config --from-literal app-color=blue: It will create a configmap with color blue

kubectl create configmap app-config --from-file /path/app.properties: It will create a config file with the app.properties content

kubectl apply -f app-config.yaml: It will create a configMap with app-config definition file

kubectl get configmaps: It will fetch all the config maps

kubectl describe configmap app-config: It will describe the config map app-config

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: nginx-configmap
data:
  name: nginx
  tier: frontend
  work: loadbalancing
```

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
  namespace: default
  labels:
    app: nginx-pod
spec:
  containers:
    - name: nginx
      image: nginx
      envFrom:
        - configMapRef:
            name: nginx-configmap
```

config map stores everything in plain text format, which is not suitable for storing password, that's why we need secrets.

We can create secret with these 3 ways.

1. By passing all the values in command.
2. By passing the properties file or yaml file in the command
3. By using the definition file.

kubectl create secret generic app-secret --from-literal app-color=blue: It will create a secret with color blue.

kubectl create secret generic app-secret --from-file app-secret.properties: It will create a secret with the app-secret.properties file.

```
#app.properties
db_host=postgres
db_password=root
db_name=database

#app-sercret.yaml
apiVersion: v1
kind: Secret
metadata:
  name: app-secret
data:
  db_host: cG9zdGdyZXMNCg==
  db_password: cm9vdA0K
  db_name: ZGF0YWJhc2UNCg==
```

```
#nginx-with-secret
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
  labels:
    name: nginx-pod
spec:
  containers:
    - name: nginx
      image: nginx
      envFrom:
        - secretRef:
            name: app-secret
```

There are some other ways to store the information. That is Helm secrets, HarshiCorp vault etc.

Security Context

In each container there can be multiple processes running and there can be multiple processes running. These processes are separated by their namespaces. By default, docker runs every process as root user, but we can change it with the following command.

Docker run --user=1000 ubuntu sleep 10

docker limits the capabilities of the root user inside the container. It is not the same root user as the host machine root user. Via **linux capabilities** we can add or remove capabilities of the root user inside the container.

Docker run --cap-add MAC_ADMIN ubuntu: It will run ubuntu container with MAC_ADMIN privilege

docker run --cap-drop kill ubuntu: It will drop the privilege of kill a process

docker run --privileged ubuntu: It will add all the privileges

This can be configured in the Kubernetes as well. In the pod level and in the container level.

```
apiVersion: v1
kind: Pod
spec:
  containers:
    - name: ubuntu-pod
      image: ubuntu
      securityContext:
        runAsUser: 1000
        capabilities:
          add: ["MAC_ADMIN", "KILL"]
```

There are 2 types of accounts in Kubernetes space:

1. User Account: Human user such as admin or developer
2. Service Account: Account created by application to interact with the Kubernetes clusters like monitoring app or build tools such as Jenkins.

Kubectl create serviceaccount dashboard-serviceaccount: It will create a service account named as dashboard-serviceaccount. It will also create a secret with a token automatically. Now with this token we can call kubernetes api. Like the following

curl <https://192.168.56.70/api> -insecure --header "Authorization Bearer #TOKEN"

We can create a service account, assign that right permission using role-based access control mechanism and expose the service account token and use it to configure the third-party application to authenticate Kubernetes api.

But in latest version we have to create the token manually and then we have to bind it with service account. **kubectl create sa dashboard-sa**

```
apiVersion: v1
kind: Secret
metadata:
  name: dashboard-sa-token
  annotations:
    kubernetes.io/service-account.name: "dashboard-sa"
type: kubernetes.io/service-account-token
```

If the third-party application is hosted in the same Kubernetes, then we can mount the service account token as volume.

We can edit the service account inside the POD, but we cannot change the service account of the deployment. After change there will be a new rollout deployment.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  selector:
    matchLabels:
      app: nginx-pod
  template:
    metadata:
      labels:
        app: nginx-pod
    spec:
      containers:
        - name: nginx
          image: nginx
          serviceAccountName: nginx-
serviceaccount
          automountServiceAccountToken: false
```

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
    - image: nginx
      name: nginx
      volumeMounts:
        - mountPath: /var/run/secrets/tokens
          name: vault-token
      serviceAccountName: build-robot
  volumes:
    - name: vault-token
      projected:
        sources:
          - serviceAccountToken:
              path: vault-token
              expirationSeconds: 7200
              audience: vault
```

By default, there is service account that is default service account.

KubectI get serviceaccount: It will fetch all the service account

kubectI describe serviceaccount default: It will describe the default service account

kubectI describe secret <token-from- sa>: With this command we can fetch the token from sa.

This default service account is associated with **image pull registry** and **image pull secrets**. We can also change it. For reference check this: <https://kubernetes.io/docs/tasks/configure-pod-container/configure-service-account/>

Resource requirements

```
apiVersion: v1
kind: LimitRange
metadata:
  name: mem-limit-range
  namespace: dev
spec:
  limits:
  - default:
      memory: "512Mi" #cpu usage
    defaultRequest:
      memory: "256Mi" #Memory usage
    type: Container
```

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
  namespace: default
  labels:
    app: nginx-pod
spec:
  containers:
  - name: nginx
    image: nginx
    resources:
      requests:
        memory: "64Mi"
        cpu: "250m"
      limits:
        memory: "128Mi"
        cpu: "500m"
```

By default, in Kubernetes there is not restriction on usage and cpu usage. With the Limit range we can set the default memory and cpu of a pod in any namespace. With resource block of the pod definition, we can also restrict the default and maximum memory usage.

KubectI apply -f mem-limit-range.yaml: It will create the memory limit of the container in any namespace

Taints and tolerations

It helps to set restrictions to what pod to schedule in which node.

Suppose we have 3 nodes and 4 pods, and we have applied a taint blue on **node-A** and we have added tolerant to **pod-B** so only **pod-B** is capable to be allocated in **node-A**. But in the same time

pod-B can also be allocated to another node. Taints imposes a rule on the node that it will only be accepting pods with specific tolerations, but it does not impose any rule on pods.

KubectI nodes <node-name> key=value:taint-effect: It will impose a taint with key value with taint effect. There are 3 taint effects 1. NoExecute 2. PrefferNoSchedule 3. NoSchedule

kubectI nodes node-a color:blue:NoSchedule

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
  namespace: default
  labels:
    app: nginx-pod
spec:
  tolerations:
    - key: "color"
      operator: "Equal"
      value: "blue"
      effect: "NoSchedule"
  containers:
    - name: nginx
      image: nginx
```

By default, the master node has some taints that prevents any pods to schedule on master.

KubectI describe node <node-name> | grep Taint: to find the taints applied on the node

kubectI nodes <node-name> key=value:taint-effect - : It will impose a taint with key value with taint effect.

KubectI get pods -o wide: To see a pod in which node

Node selectors and Node Affinity

kubectI label node <node-name> key=value: To label a node

The specific pod with specific node with **nodeSelector with key value** will be placed on that node.

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
  namespace: default
  labels:
    app: nginx-pod
spec:
  nodeSelector:
    app: color
  containers:
    - name: nginx
      image: nginx
```

Node selector is easy to use but lack the advanced features that is why we are using node Affinity. Here also at first, we must label the node

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
  namespace: default
  labels:
    app: nginx-pod
spec:
  affinity:
    nodeAffinity:
      requiredDuringSchedulingIgnoredDuringExecution:
        nodeSelectorTerms:
          - matchExpressions:
              - key: size
                operator: In
                values:
                  - "Large"
                  - "Medium"
  containers:
    - name: nginx
      image: nginx
```

To place this nginx pod we must label the node with Large or Medium else pod will not be executed in the node.

There are 3 types of operators:

1. **In:** If the any value from values is there in the labels in the node
2. **No:** If the values from the values are there not in labels of thr node
3. **Exists:** If the key exists in the label of the node

There are 3 types of nodeAffinity selectors.

1. `requiredDuringSchedulingIgnoredDuringExecution`
2. `prefferedDuringSchedulingIgnoredDuringExecution`
3. `requiredDuringSchedulingRequiredDuringExecution` (Available in future)

Name	During scheduling	During execution
<code>requiredDuringSchedulingIgnoredDuringExecution</code>	Label is required in the node other wise pod will not be executed	In if there is any change in label in node then pod will ignore that and continue executing
<code>prefferedDuringSchedulingIgnoredDuringExecution</code>	It will try to find the node with the label. If it does not find the node, then it tries to place the pod in any node	In if there is any change in label in node then pod will ignore that and continue executing
<code>requiredDuringSchedulingRequiredDuringExecution</code>	Label is required in the node otherwise pod will not be executed	In if there is any change in label in node then pod should have that label other wise it will stop executing

With the combination of taints and node affinity we can specify that which pod will place on which node.

Multi container pod

A multi container is something that has more than one container in pod.

1. **Side car** -> Log agent with a web server
2. **Adapter** -> Log agent that process different type of web server and pushes in the central log server.
3. **Ambassador**-> Microservice with DB agent which connects to different type and environments of database.

Readiness and liveness probe

Readiness probe: Sometimes container is up that does not mean that the container is ready for external communication like Jenkins takes time to boot up. So, we can configure an api or a script to check the container is ready or not.

Liveness probe: It is way to check periodically that the application is healthy or not, otherwise it destroys the container and starts a new one.

There are 3 ways to check:

1. httpGet
2. tcpSocket
3. start-up script

we can also add `initialDelaySeconds`, `periodSeconds`, `failureThreshold` configure other options

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
  namespace: default
  labels:
    app: nginx-pod
spec:
  containers:
    - name: nginx
      image: nginx
      readinessProbe:
        httpGet:
          path: /api/ready
          port: 8080
          httpHeaders:
            - name: Custom-Header
              value: Awesome
          initialDelaySeconds: 10
          periodSeconds: 5
          failureThreshold: 8
      livenessProbe:
        httpGet:
          path: /api/ready
          port: 8080
          httpHeaders:
            - name: Custom-Header
              value: Awesome
          initialDelaySeconds: 10
          periodSeconds: 5
          failureThreshold: 8
```

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
  namespace: default
  labels:
    app: nginx-pod
spec:
  containers:
    - name: nginx
      image: nginx
      readinessProbe:
        tcpSocket:
          port: 8080
      livenessProbe:
        tcpSocket:
```

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
  namespace: default
  labels:
    app: nginx-pod
spec:
  containers:
    - name: nginx
      image: nginx
      readinessProbe:
        exec:
          - "cat"
          - "/app/ready"
      livenessProbe:
        exec:
          - "cat"
          - "/app/ready"
```


Container logging

docker run -d kodecloud/event-simulator

docker log -f <container-name>

this is how we can check the logs of a docker container after running it in detached mode. We can do the same with Kubernetes.

kubectrl create -f event-simulator.yaml

kubectrl logs -f event-simulator-pod: To get logs of a specific pod

kubectrl logs -f <pod-name> <container-name>: To get the logs of a specific container of a specific pod

Monitoring Solutions:

1. Prometheus
2. Elastic stack
3. Datadog
4. Dynatree

Kubernetes runs an agent on each node known as kubelet which is responsible for receiving instruction from Kubernetes master server. It has a subcomponent known as container advisor, it is responsible for retrieving performance metrics from pods and exposing them to kubelet api.

Minikube addons enable metrics-server: To enable the metric server on minikube

kubectrl top node: To find the node with max using the resources

kubectrl top pod: To find the pod with the max using resources.

Pod Design

Labels and selectors are standard methods to group things together. We can keep these labels under metadata. These labels are used to identify a specific pod by the service, replica set and deployment. Though one label is sufficient to identify but we can use as many labels as we want.

Alongside with labels, we can also add annotations. It used to save build information and other things to documentation purpose.

kubectrl get pods -l env=dev: It will find the pods where label is dev.

kubectrl get pods -l env=dev --no-headers | wc -l: It will find the count of the pods with this label.

kubectrl get all -l env=prod: It will find all the objects in prod environment.

Kubectrl get all -l env=prod --no-headers | wc -l : It will count all the objects in prod environment.

Kubectl get pods -l env=prod,bu=finance,tier=frontend: It will find all the pods with all these label.

When we first create a deployment, it creates a rollout, and a new rollout creates a new deployment revision. When there is a change in deployment like image or something then again it creates a rollout which again creates a revision. It helps us to keep track of the changes.

There are two types of deployment strategy.

1. **Recreate:** Destroy all the container then again create the container with changes. During the period when the older version is down the application become inaccessible.
2. **Rolling update:** In this strategy some of the old containers are down and int that place new containers with change are started. If there is no mentioning of the deployment strategy in the yaml file, then this strategy become the default strategy.

If we perform the **kubectl describe deployment <deployment-name>** then in the events, we can see the events how containers went down and up.

When there is upgrade the deployment object internally creates a replicate set and then fill that with the new containers and delete the containers from old replica set.

To undo a change, use the following command

kubectl rollout undo deployment/<deployment-name>

In rollback also deployment object creates another replica set and fill that.

Summarize commands:

kubectl create -f <deployment-definition.yml>

kubectl get deployments

kubectl apply -f <deployment-definition.yml>

kubectl set image deployment/<deployment-name> nginx=nginx:1.9.1

kubectl rollout status deployment/<deployment-name>

kubectl rollout history deployment/<deployment-name>

kubectl rollout undo deployment/<deployment-name>

Updating a Deployment

Here are some handy examples related to updating a Kubernetes Deployment:

Creating a deployment, checking the rollout status and history:

In the example below, we will first create a simple deployment and inspect the rollout status and the rollout history:

```
master $ kubectl create deployment nginx --image=nginx:1.16
deployment.apps/nginx created

master $ kubectl rollout status deployment nginx
Waiting for deployment "nginx" rollout to finish: 0 of 1 updated replica are available...
deployment "nginx" successfully rolled out

master $ kubectl rollout history deployment nginx
deployment.extensions/nginx
REVISION CHANGE-CAUSE
1    <none>
```

Using the `--revision` flag:

Here the revision 1 is the first version where the deployment was created.

You can check the status of each revision individually by using the `--revision` flag:

```
master $ kubectl rollout history deployment nginx --revision=1
deployment.extensions/nginx with revision #1

Pod Template:
Labels: app=nginx pod-template-hash=6454457cdb
Containers: nginx: Image: nginx:1.16
Port: <none>
Host Port: <none>
Environment: <none>
Mounts: <none>
Volumes: <none>
master $
```

Using the `--record` flag:

You would have noticed that the “change-cause” field is empty in the rollout history output. We can use the `--record` flag to save the command used to create/update a deployment against the revision number.

You can now see that the change-cause is recorded for the revision 2 of this deployment.

Let’s make some more changes. In the example below, we are editing the deployment and changing the image from `nginx:1.17` to `nginx:latest` while making use of the `--record` flag.

```
master $ kubectl set image deployment nginx nginx=nginx:1.17 --record
deployment.extensions/nginx image updated
master $master $
```

```
master $ kubectl rollout history deployment nginx
deployment.extensions/nginx
```

REVISION CHANGE-CAUSE

- 1 <none>
- 2 kubectl set image deployment nginx nginx=nginx:1.17 --record=true

```
master $ kubectl edit deployments. nginx --record
deployment.extensions/nginx edited
```

```
master $ kubectl rollout history deployment nginx
```

REVISION CHANGE-CAUSE

- 1 <none>
- 2 kubectl set image deployment nginx nginx=nginx:1.17 --record=true
- 3 kubectl edit deployments. nginx --record=true

```
master $ kubectl rollout history deployment nginx --revision=3
deployment.extensions/nginx with revision #3
```

Pod Template: Labels: app=nginx

pod-template-hash=df6487dc Annotations: kubernetes.io/change-cause: kubectl edit
deployments. nginx --record=true

Containers:

nginx:

Image: nginx:latest

Port: <none>

Host Port: <none>

Environment: <none>

Mounts: <none>

Volumes: <none>

Undo a change:

Let's now rollback to the previous revision:

```
master $ kubectl rollout undo deployment nginx
deployment.extensions/nginx rolled back

master $ kubectl rollout history deployment nginx
deployment.extensions/nginxREVISION CHANGE-CAUSE
1  <none>
3  kubectl edit deployments. nginx --record=true
4  kubectl set image deployment nginx nginx=nginx:1.17 --record=true
```

```
master $ kubectl rollout history deployment nginx --revision=4
deployment.extensions/nginx with revision #4Pod Template:
Labels:  app=nginx  pod-template-hash=b99b98f9
Annotations:  kubernetes.io/change-cause: kubectl set image deployment nginx
nginx=nginx:1.17 --record=true
Containers:
  nginx:
    Image:  nginx:1.17
    Port:  <none>
    Host Port: <none>
    Environment:  <none>
    Mounts:  <none>
    Volumes:  <none>

master $ kubectl describe deployments. nginx | grep -i image:
Image:  nginx:1.17
```

With this, we have rolled back to the previous version of the deployment with the image = nginx:1.17.

Jobs and CronJobs

When we have a requirement that we need a certain type of job to execute and finish like batch processing, data processing, analytics then on that time rather than creating a pod we can create any job.

Because the pods keep on restarting when it got terminated but we need to start, execute and finish. The task which are meant to live for a short period of time.

Like when we do the following command:

docker run ubuntu expr 3+2

It will run execute and terminate, since the task was completed.

We can replicate the same with Kubernetes with job.

```
apiVersion: v1
kind: Pod
metadata:
  name: Math-app-pod
  labels:
    name: math-app-pod
spec:
  containers:
  - name: math-app-pod
    image: ubuntu
    command: ["expr", "3", "+", "2"]
```

By default, ubuntu will run execute and finish and so did the pod. It will go to the complete state.

But in attempt of keep on running Kubernetes will keep this pod restarting. This will happen continuously until the threshold is reached. This is default behaviour of the pod and there is a property mentioned under the **spec** which is **restartPolicy**. By default, its value is **Always**.

We can set this value to **Never** to stop restarting.

We can create a job in place of this pod.

```
apiVersion: batch/v1
kind: Job
metadata:
  name: math-app-job
  labels:
    name: math-app-job
spec:
  template:
    spec:
      containers:
      - name: math-app-pod
        image: ubuntu
        command: ["expr", "3", "+", "2"]
        restartPolicy: Never
```

kubectl create -f math-add-job.yaml: To create a job

kubectl get jobs: To get all the jobs

Kubernetes internally create a pod for this and once it got completed it will not try to restart.

We can see the output from the logs of the pod.

kubectl logs <pod-name>

To run multiple jobs, we can add **completion** property under **spec**. By default, the pods are created one after another like 2nd pod is created only after the 1st pod is started and finished. If the pod got failed it will try to bring a new pod until it reaches the required number of jobs mentioned. So, we can run the jobs parallelly to save some time. To add parallelism, we can add **parallelism** property under the **spec**.

```
spec:
  completions: 3
  parallelism: 3
  template:
    spec:
      containers:
      - name: random-error-job
        image: kodekloud/random-error
        restartPolicy: Never
```

CronJob is a job that can be scheduled with cron expression

```
apiVersion: batch/v1beta1
kind: CronJob
metadata:
  name: math-cronjob
  labels:
    name: math-cronjob
spec:
  schedule: "*/1 * * * *"
  jobTemplate:
    spec:
      completions: 3
      parallelism: 3
      template:
        spec:
          containers:
          - name: math-cronjob
            image: ubuntu
            command: ["expr", "3", "+", "2"]
            restartPolicy: Never
```

kubectl create -f math-cronjob.yaml

kubectl get cronjob

Kubernetes Services/Ingress

Some references for ingress:

<https://medium.com/digitalfrontiers/kubernetes-ingress-with-nginx-93bdc1ce5fa9>

<https://medium.com/devops-mojo/kubernetes-ingress-overview-what-is-kubernetes-ingress-introduction-to-k8s-ingress-b0f81525ffe2>

<https://katharharshal1.medium.com/setup-nginx-ingress-controller-on-kubernetes-cluster-dd48b2b1ab61>

<https://www.containiq.com/post/kubernetes-ingress>

<https://www.edc4it.com/blog/k8s-ingress-tls-termination>

ingress controller:

<https://raw.githubusercontent.com/kubernetes/ingress-nginx/nginx-0.27.0/deploy/static/mandatory.yaml>

<https://raw.githubusercontent.com/kubernetes/ingress-nginx/nginx-0.27.0/deploy/static/provider/cloud-generic.yaml>

minikube addons enable ingress

minikube addons enable metrics-server

minikube dashboard

kubectrl proxy --port=8080

curl http://localhost:8080/api/

<https://github.com/nginxinc/kubernetes-ingress>

<https://kubernetes.io/docs/concepts/services-networking/ingress/>

kubectrl get all --all-namespaces

kubectrl get cm --all-namespaces

kubectrl describe cm kube-root-ca.crt -n <namespace>

kubectrl delete services,deployments,rs,pods,cm,secrets,job.batch --all -n <namespace>

kubectrl get cm,secrets --all-namespaces

kubectrl delete --all -n ingress-nginx

kubectrl create deployment hello-minikube --image=kicbase/echo-server:1.0

kubectrl expose deployment hello-minikube --type=NodePort --port=8080

kubectrl get services hello-minikube

minikube service hello-minikube

kubectrl port-forward service/hello-minikube 7080:8080

kubectrl create deployment balanced --image=kicbase/echo-server:1.0


```
kubectl expose deployment balanced --type=LoadBalancer --port=8080
minikube tunnel
kubectl get services balanced
```

```
minikube config set memory 9001
minikube addons list
minikube start -p aged --kubernetes-version=v1.16.1
minikube delete --all
```

Now, in k8s version **1.20+** we can create an Ingress resource from the imperative way like this:-

Format - kubectl create ingress <ingress-name> --rule="host/path=service:port"

Example - kubectl create ingress ingress-test --rule="wear.my-online-store.com/wear*=wear-service:80"

Find more information and examples in the below reference link:-

<https://kubernetes.io/docs/reference/generated/kubect/kubectl-commands#-em-ingress-em->

References:-

<https://kubernetes.io/docs/concepts/services-networking/ingress>

<https://kubernetes.io/docs/concepts/services-networking/ingress/#path-types>

if minikube driver is none then follow this:

minikube addons enable ingress / else download all the ingress controller configs and apply them

```
kubectl apply -f namespace.yaml
kubectl apply -f deployments.yaml
kubectl apply -f services.yaml
kubectl apply -f ingress.yaml / ingress-resource.yaml
```

minikube tunnel

update /etc/host file -> 127.0.0.1 app.reactor.com

we are setting 127.0.0.1 not 192.168.49.2 because in minikube local ingress is in 127.0.0.1

```
14:35:14:~ kubectl get ingress -A
NAMESPACE NAME CLASS HOSTS ADDRESS PORTS AGE
default echo <none> app.reactor.com 192.168.49.2 80 6h6m
```

Don't use this annotation in metadata `nginx.ingress.kubernetes.io/rewrite-target: /`

It is used to rewrite the path. `/echo/sample` will be converted to `/`

Ingress resource port is always set to 80/443

```

---
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: echo
  # annotations:
  #   nginx.ingress.kubernetes.io/rewrite-
target: /
spec:
  rules:
    - host: app.reactor.com
      http:
        paths:
          - backend:
              service:
                name: echo
              port:
                number: 80
            pathType: Prefix
            path: /echo

```

```

---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: echo
  labels:
    type: echo
spec:
  replicas: 2
  selector:
    matchLabels:
      type: echo
  template:
    metadata:
      labels:
        type: echo
    spec:
      containers:
        - name: echo
          image: abhishek1009/echo:latest
          ports:
            - containerPort: 8000
---
apiVersion: v1
kind: Service
metadata:
  name: echo
spec:
  selector:
    type: echo
  ports:
    - port: 80
      targetPort: 8000

```

We can write as many rules as we want.

Network Policies

Let's say there is a DB pod and one frontend pod and one backend pod. If we create a service to expose the DB pod via service into some specific port then the DB pod will be accessible by all other pods, but we don't want that to happen. So we can create some rules or **NetworkPolicy** on DB pod that who can access the pod and whom the db pod can access. There are the ingress policy and egress policy. As name suggest ingress means incoming traffic rule to specific pod and egress means outgoing traffic rules to specific pods.

Kubernetes is configured with all allow rule, that means any pod in any node can connect to any pod. To apply any ingress or egress rule we have to apply **NetworkPolicy**. Just like services and replicaset/deployment Kubernetes will use label and selectors to apply the policies.

```
---
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: db-policy
spec:
  podSelector:
    matchLabels:
      type: master-db
  policyTypes:
    - Ingress
  ingress:
    - from:
        - podSelector:
            matchLabels:
              type: backend-api
          namespaceSelector:
            matchLabels:
              env: prod
        - ipBlock:
            cidr: 192.168.5.10/32
      ports:
        - protocol: TCP
          port: 3306
```

```
---
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: nginx-policy
spec:
  podSelector:
    matchLabels:
      app: nginx
  policyTypes:
    - Ingress
    - Egress
  ingress:
    - from:
        - podSelector:
            matchLabels:
              type: frontend-api
          namespaceSelector:
            matchLabels:
              env: prod
        - ipBlock:
            cidr: 192.168.5.10/32
      ports:
        - protocol: TCP
          port: 80
  egress:
    - to:
        - podSelector:
            matchLabels:
              type: backend-api
          namespaceSelector:
            matchLabels:
              env: prod
        - ipBlock:
            cidr: 192.168.5.17/32
      ports:
        - protocol: TCP
          port: 8080
```